

MMAI 5500 Assignment 2

DIY candy counter

In this assignment you will teach an object detector to count different types of candy. You will annotated images of candy and fine-tune an object detector adapting a [Huggingface tutorial](#). The tutorial will have to be modified to work with your dataset (see below).

Data

Data labelling

You will label your own dataset using [Label Studio](#). Install it on your local machine by following the instructions on Label Studio [Quick Start](#).

Label the 10 images in `candy_images.zip` with the following 8 candy types:

```
{
  'Moon': 1,
  'Insect': 2,
  'Black_star': 3,
  'Grey_star': 4,
  'Unicorn_whole': 5,
  'Unicorn_head': 6,
  'Owl': 7,
  'Cat': 8
}
```

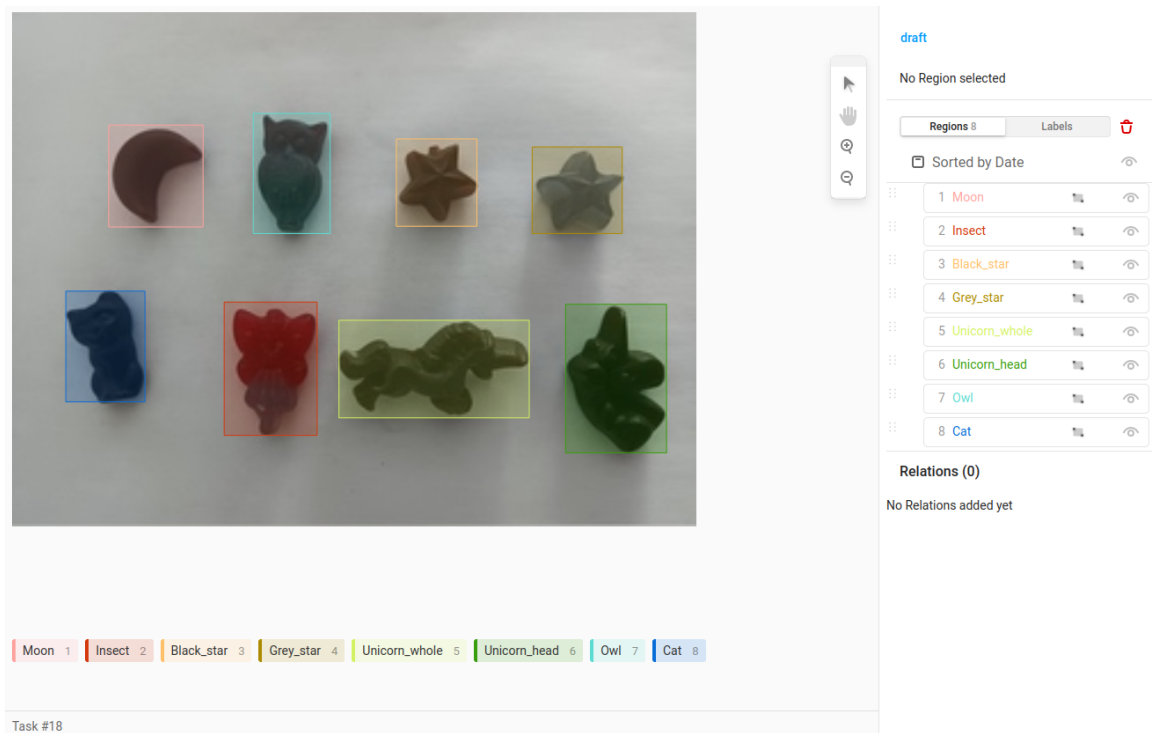
More labelling examples are included in the ZIP file `MMAI5500_Assignment2.zip`. Make the bounding boxes as tight around the objects as possible.

Export the images and annotations in the COCO json format. The export should return a ZIP archive containing a folder `images` and a file `result.json`. The former contains the images and the latter is a `json` file with the annotations (i.e. bounding boxes and class labels).

Load the data with Huggingface's DatasetDict

The annotated data exported from Label Studio cannot be directly loaded with Huggingface's `datasets.load_dataset()`. Below are some code snippets to help with the conversion.

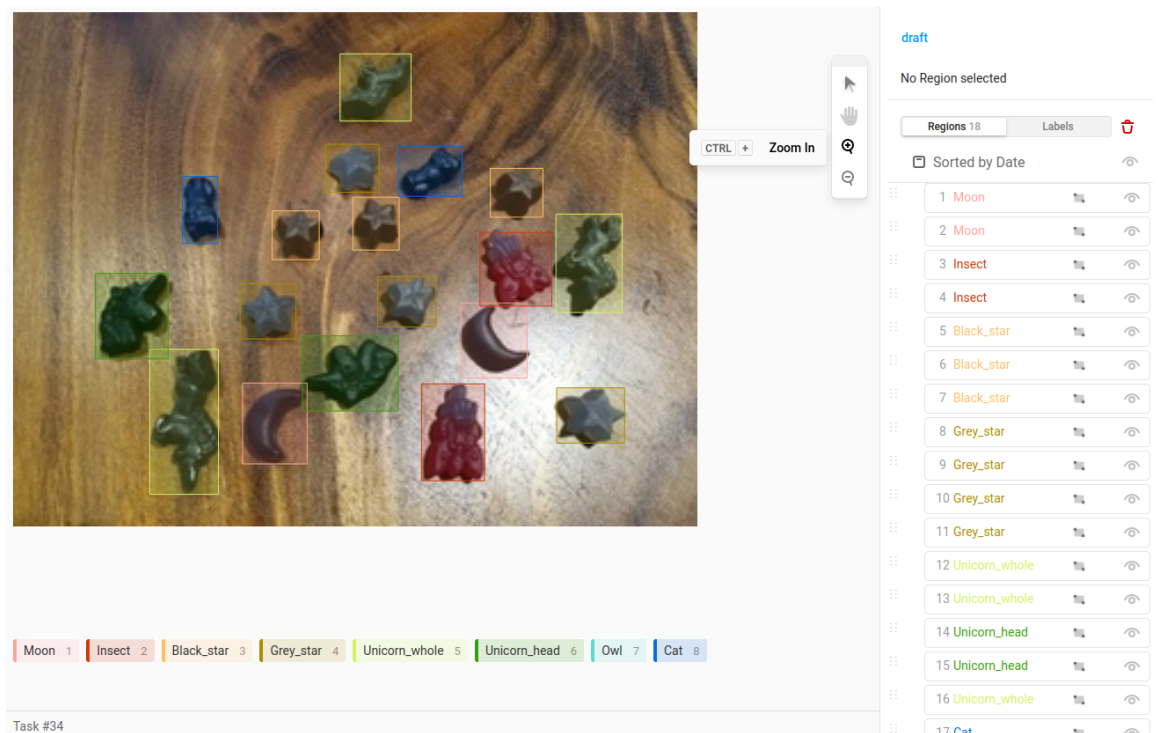
1. Read the COCO formatted annotations to a dict



```
import json
with open('result.json') as f:
    cocodata = json.load(f)
```

2. Convert the COCO formatted dict to a format readable by `dataset.load_dataset()`.

```
import os
# Store Huggingface formatted data in a list
huggingdata = []
# Iterate through the images
for image in cocodata['images']:
    # Remove the image directory from the file name
    image['file_name'] = image['file_name'].split(os.path.sep)[-1]
    image['image_id'] = image['id']
    # Extend the image dict with bounding boxes and class labels
    image['objects'] = {'bbox': [], 'category': [], 'area': [], 'id': []}
    # Iterate through the annotations (bounding boxes and labels)
    for annot in cocodata['annotations']:
        # Check if the annotation matches the image
        if annot['image_id'] == image['id']:
            # Add the annotation
```



```

image['objects']['bbox'].append(annot['bbox'])
image['objects']['category'].append(annot['category_id'])
image['objects']['area'].append(annot['area'])
image['objects']['id'].append(annot['id'])
# Append the image dict with annotations to the list
huggingdata.append(image)

```

3. Write the Huggingface formatted annotations to a json file.

```

with open("metadata.jsonl", 'w') as f:
    for item in huggingdata:
        f.write(json.dumps(item) + "\n")

```

4. Follow [Create an image dataset](#) and the section [Create an image dataset: object detection](#) to organize the images and metadata.jsonl into the correct folder structure.
5. Read the data into a [DatasetDict](#).

```

from datasets import load_dataset
# Assumes the data is stored in a folder called "data".
candy_data = load_dataset('imagefolder', data_dir="data")

```

To convert from ids to labels and back again (labels to ids) you can use this snippet:

```
id2label = {item['id']: item['name'] for item in cocodata['categories']}  
label2id = {v: k for k, v in id2label.items()}
```

Save the model

Save the model using `trainer.save_model('candy_detector')` and download it from Colab.

Tips and tricks

Training

Your dataset is very small and much smaller than the `cppe5` dataset used in the tutorial. Thus, you will probably have to finetune your model for more epochs than in the tutorial. You can set the number of epochs in `TrainingArguments()` with the argument `num_train_epochs`. Make sure that the models learns, i.e. you see the training loss decreasing. `trainer.train()` will print the training loss according the the `logging_steps` set in the `TrainingArguments()`. If you set the number of `logging_steps=10` then the loss will be printed every 10 steps (batches).

Detection

Test your model by running inference/detection on some images. When you do this, you might not get any detections. A possible reason for this is that you set the detection threshold too high. In the tutorial they use `threshold=0.5` but this might be too high for your model. Try lowering the threshold.

With the `pipeline` you can set the threshold as follows:

```
obj_detector = pipeline("object-detection", model="candy_detector")  
obj_detector(image, threshold=0.5)
```

Deliverables

You need to submit a notebook (IPYNB) and the trained model. The notebook should prepare and load the data, fine-tune the model, evaluate and save it. It should also have a method called `candy_counter(image)` that loads the fine-tuned model that takes as input a picture of candy (a Numpy array with the

same shape as the original images) and returns a dictionary with the counts of the different types of candies. `candy_counter(image)` should use the [pipeline](#) for inference. All the cells should be correctly executed, including a test image with predicted boundary boxes and saved in the notebook.

Example output:

```
In [6]: candy_counter(image)
Out[6]:
{'Moon': 5,
 'Insect': 0,
 'Black_star': 0,
 'Grey_star': 6,
 'Unicorn_whole': 2,
 'Unicorn_head': 1,
 'Owl': 0,
 'Cat': 4}
```

Grading

For full marks the submitted code needs to be bug free, include everything described under **Deliverables**, and the `candy_counter` has to have a [mAP](#) score above 0.5 on unseen test data.

Good luck!